


mongoDB

Exploring the Aggregation Framework

Agenda

1. Analytics in MongoDB?
2. Aggregation Framework
3. Aggregation Framework in Action
 - US Census Data
1. Aggregation Framework Options

Analytics in MongoDB?



For Example: US Census Data

- Census data from 1990, 2000, 2010
- Question:
 - Which US Division has the fastest growing population density?
 - We only want to include data for states with more than 1M people
 - We only want to include divisions larger than 100K square miles

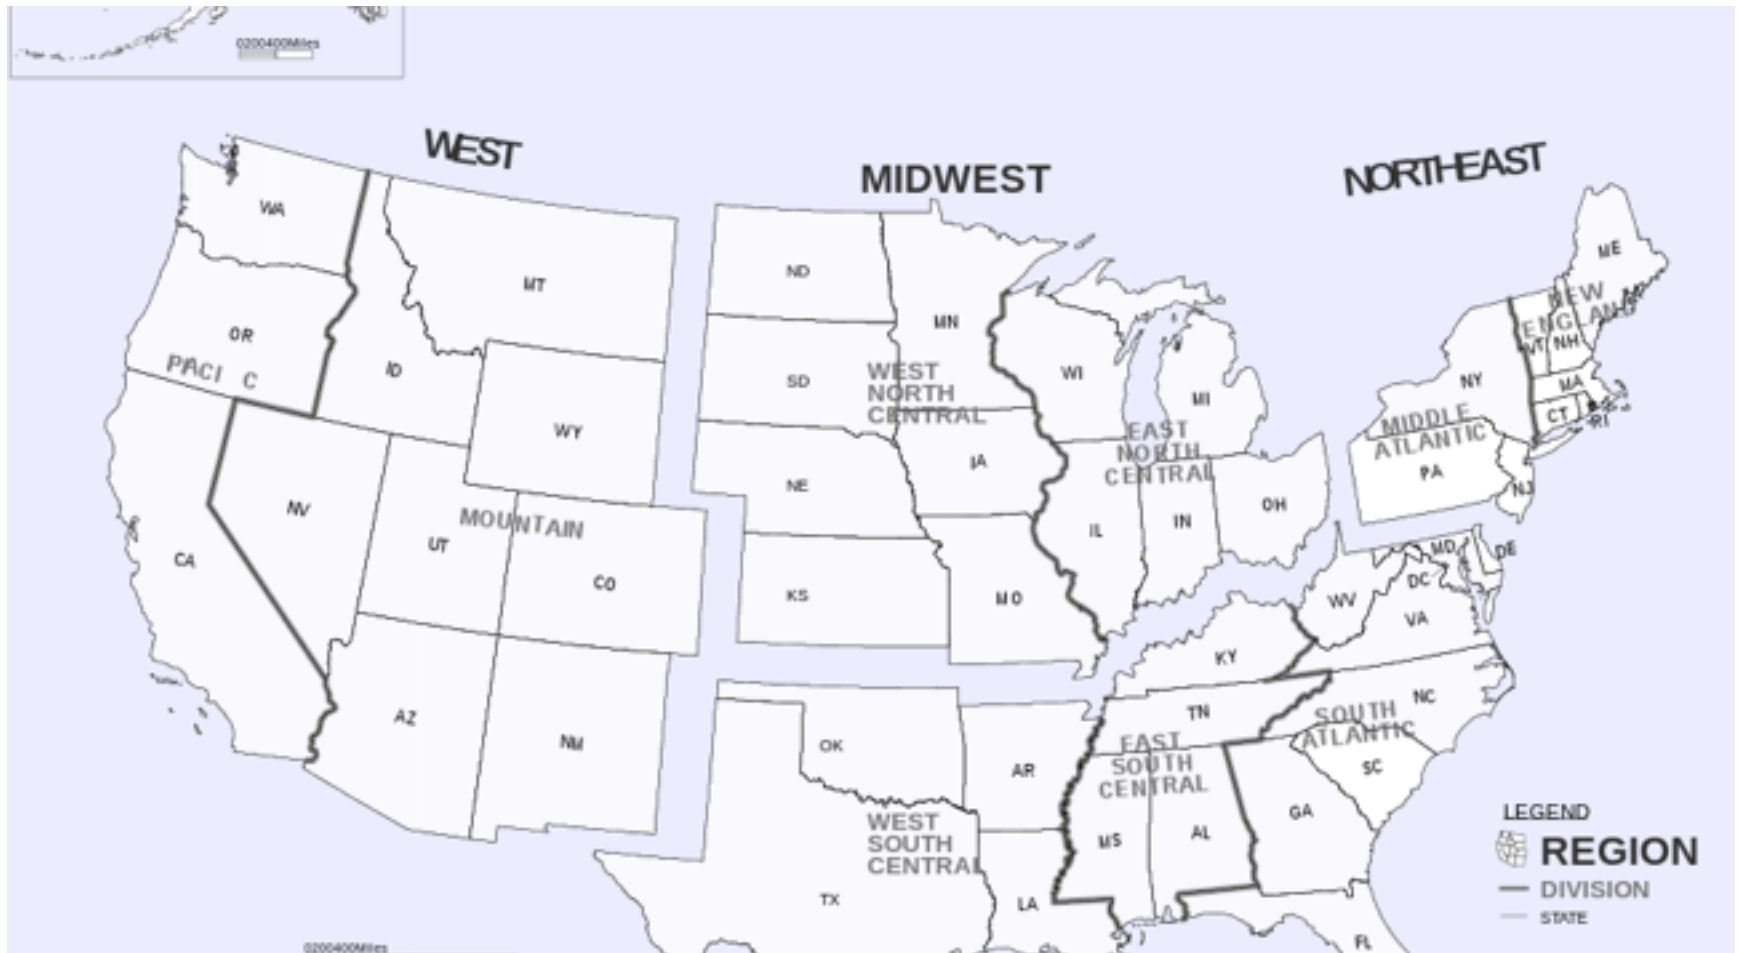


Division = a group of US States

Population density = Area of division/# of people

Data is provided at the state level

US Regions and Divisions



How would we solve this in SQL?

- SELECT GROUP BY HAVING

What About MongoDB?

The Aggregation Framework

What is an Aggregation Pipeline?

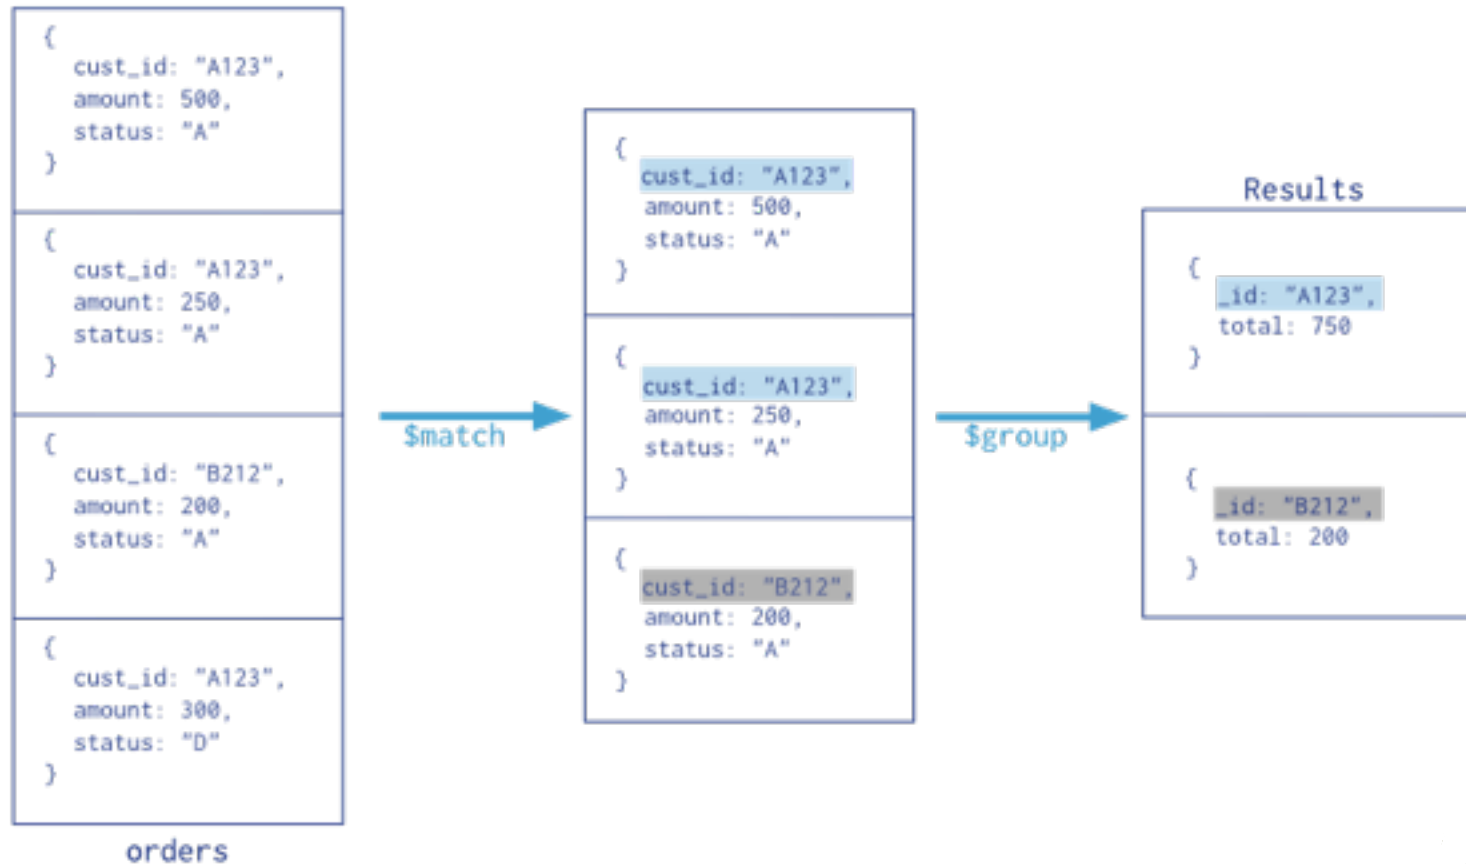
- A Series of Document Transformations
 - Executed in stages
 - Original input is a collection
 - Output as a cursor or a collection



- Rich Library of Functions
 - Filter, compute, group, and summarize data
 - Output of one stage sent to input of next
 - Operations executed in sequential order

Aggregation Pipeline

Collection
↓
`db.orders.aggregate([`
 `$match phase → { $match: { status: "A" } },`
 `$group phase → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]`)



Pipeline Operators

- `$match`
Filter documents
- `$project`
Reshape documents
- `$group`
Summarize documents
- `$unwind`
Expand documents
- `$sort`
Order documents
- `$limit/$skip`
Paginate documents
- `$redact`
Restrict documents
- `$geoNear`
Proximity sort documents
- `$let, $map`
Define variables

Aggregation Framework in Action

(let's play with the census data)



MongoDB State Collection

- Document For Each State
 - Name
 - Region
 - Division
- Census Data For 1990, 2000, 2010
 - Population
 - Housing Units
 - Occupied Housing Units
- Census Data is an array with three subdocuments

Document Model

```
{  "_id" : ObjectId("54e23c7b28099359f5661525"),
    "name" : "California",
    "region" : "West",
    "data" : [
      { "totalPop" : 33871648,
        "totalHouse" : 12214549,
        "occHouse" : 11502870,
        "year" : 2000 },
      { "totalPop" : 37253956,
        "totalHouse" : 13680081,
        "occHouse" : 12577498,
        "year" : 2010 },
      { "totalPop" : 29760021,
        "totalHouse" : 11182882,
        "occHouse" : 29008161,
        "year" : 1990 }
    ],
    ...
}
```

Total US Area

```
db.cData.aggregate([  
  {"$group" : {"_id" : null,  
    "totalArea" : {$sum : "$areaM"},  
    "avgArea" : {$avg : "$areaM"}}}] )
```

\$group

- Group documents by value
 - Field reference, document, constant
 - Other output fields are computed
 - `$max`, `$min`, `$avg`, `$sum`
 - `$addToSet`, `$push`
 - `$first`, `$last`
 - Processes all data in memory by default



Area By Region

```
db.cData.aggregate([
  {"$group" : {"_id" : "$region",
    "totalArea" : {$sum : "$areaM"},
    "avgArea" : {$avg : "$areaM"},
    "numStates" : {$sum : 1},
    "states" : {$push : "$name"}}}]
])
```

Calculating Average State Area By Region

```
{  
  state: "New York",  
  areaM: 54554,  
  region: "Northeast"  
}
```



```
{ $group: {  
  _id: "$region",  
  avgArea: { $avg: "$areaM" }  
}}
```



```
{  
  state: "New Jersey",  
  areaM: 8722,  
  region: "Northeast"  
}
```

```
{  
  _id: "Northeast",  
  avgAreaM: 20146  
}
```

```
{  
  state: "California",  
  areaM: 163694,  
  region: "West"  
}
```

```
{  
  _id: "West",  
  avgAreaM: 144096  
}
```

Calculating Total Area and State Count

```
{
  state: "New York",
  areaM: 54554,
  region: "North East"
}
```



```
{ $group: {
  _id: "$region",
  totArea: {$sum:
    "$areaM" },
  sCount : {$sum : 1}}}
```



```
{
  state: "New Jersey",
  areaM: 8722,
  region: "North East"
}
```

```
{
  _id: "Northeast",
  totArea: 308
  sCount: 2}
```

```
{
  state: "California",
  area: 163694,
  region: "West"
}
```

```
{
  _id: "West",
  totArea: 300,
  sCount: 1}
```

Total US Population By Year

```
db.cData.aggregate(  
  [{ $unwind : "$data" },  
    { $group : { "_id" : "$data.year",  
                  "totalPop" : { $sum : "$data.totalPop" } } },  
    { $sort : { "totalPop" : 1 } }  
  ] )
```

\$unwind

- Operate on an array field
 - Create documents from array elements
 - Array replaced by element value
 - Missing/empty fields → no output
 - Non-array fields → error
 - Pipe to \$group to aggregate

\$unwind

```
{  
  state: "New York",  
  census: [1990, 2000,  
           2010]  
}
```



```
{ $unwind: $census }
```



```
{  
  state: "New Jersey",  
  census: [1990, 2000]  
}
```

```
{ state: "New York",  
  census: 1990 }
```

```
{ state: "New York",  
  census: 2000 }
```

```
{  
  state: "California",  
  census: [1980, 1990,  
           2000, 2010]  
}
```

```
{ state: "New York",  
  census: 2010 }
```

```
{ state: "New Jersey",  
  census: 1990 }
```

```
{  
  state: "Delaware",  
  census: [1990, 2000]  
}
```

```
{ state: "New Jersey",  
  census: 2000 }
```

⋮

Southern State Population By Year

```
db.cData.aggregate(  
  [{ $match : { "region" : "South" } },  
  { $unwind : "$data" },  
  { $group : { "_id" : "$data.year"  
               "totalPop" : { "$sum"  
                             "$data.totalPop" } } } ] )
```

Include example with compound group?

\$match

- Filter documents
 - Uses existing query syntax
 - No \$where (server side Javascript)



\$match

```
{  
  state: "New York",  
  areaM: 218,  
  region: "Northeast"  
}
```



```
{ $match:  
  { "region" : "West" }  
}
```



```
{  
  state: "Oregon",  
  areaM: 245,  
  region: "West"  
}
```

```
{  
  state: "Oregon",  
  areaM: 245,  
  region: "West"  
}
```

```
{  
  state: "California",  
  area: 300,  
  region: "West"  
}
```

```
{  
  state: "California",  
  area: 300,  
  region: "West"  
}
```

Population Delta By State from 1990 to 2010


```
db.cData.aggregate(  
  [{ $unwind : "$data" },  
    { $sort : { "data.year" : 1 } },  
    { $group : { "_id" : "$name",  
                  "pop1990" : { "$first" : "$data.totalPop" },  
                  "pop2010" : { "$last" : "$data.totalPop" } } },  
    { $project : { "_id" : 0,  
                  "name" : "$_id",  
                  "delta" : { "$subtract" :  
                                [ "$pop2010", "$pop1990" ] },  
                  "pop1990" : 1,  
                  "pop2010" : 1 }  
    } ]  
)
```

Population Delta By State from 1990 to 2010

```
db.cData.aggregate(  
  [{ $unwind : "$data" },  
    { $sort : { "data.year" : 1 } },  
  { $group : { "_id" : "$name",  
    "pop1990" : { "$first" : "$data.totalPop" },  
    "pop2010" : { "$last" : "$data.totalPop" } } },  
  { $project : { "_id" : 0,  
    "name" : "$_id",  
    "delta" : { "$subtract" :  
      [ "$pop2010", "$pop1990" ] },  
    "pop1990" : 1,  
    "pop2010" : 1 }  
  } ]  
)
```

\$sort, \$limit, \$skip

- Sort documents by one or more fields
 - Same order syntax as cursors
 - Waits for earlier pipeline operator to return
 - In-memory unless early and indexed



- Limit and skip follow cursor behavior

Population Delta By State from 1990 to 2010

```
db.cData.aggregate(  
  [{ $unwind : "$data" },  
    { $sort : { "data.year" : 1 } },  
    { $group : { "_id" : "$name",  
      "pop1990" : { "$first" : "$data.totalPop" },  
      "pop2010" : { "$last" : "$data.totalPop" } } },  
    { $project : { "_id" : 0,  
      "name" : "$_id",  
      "delta" : { "$subtract" :  
        [ "$pop2010", "$pop1990" ] },  
      "pop1990" : 1,  
      "pop2010" : 1 }  
    } ]  
)
```

\$first, \$last


- Collection operations like \$push and \$addToSet
- Must be used in \$group
- \$first and \$last determined by document order
- Typically used with \$sort to ensure ordering is known

Population Delta By State from 1990 to 2010

```
db.cData.aggregate(  
  [{ $unwind : "$data" },  
    { $sort : { "data.year" : 1 } },  
    { $group : { "_id" : "$name",  
      "pop1990" : { "$first" : "$data.totalPop" },  
      "pop2010" : { "$last" : "$data.totalPop" } } },  
    { $project : { "_id" : 0,  
      "name" : "$_id",  
      "delta" : { "$subtract" :  
        [ "$pop2010", "$pop1990" ] },  
      "pop1990" : 1,  
      "pop2010" : 1 }  
    } ]  
)
```

\$project

- Reshape Documents
 - Include, exclude or rename fields
 - Compute field values
 - Create sub-document fields



Including and Excluding Fields

```
{
  "_id" : "Virginia",
  "pop1990" : 453588,
  "pop2010" : 3725789
}
```



```
{ $project:
  { "_id" : 0,
    "pop1990" : 1,
    "pop2010" : 1
  }
}
```



```
{
  "_id" : "South Dakota",
  "pop1990" : 453588,
  "pop2010" : 3725789
}
```

```
{
  "pop1990" : 453588,
  "pop2010" : 3725789
}
```

```
{
  "pop1990" : 453588,
  "pop2010" : 3725789
}
```

Renaming and Computing Fields

```
{
  "_id" : "Virginia",
  "pop1990" : 6187358,
  "pop2010" : 8001024
}
```



```
{
  "_id" : "South Dakota",
  "pop1990" : 696004,
  "pop2010" : 814180
}
```

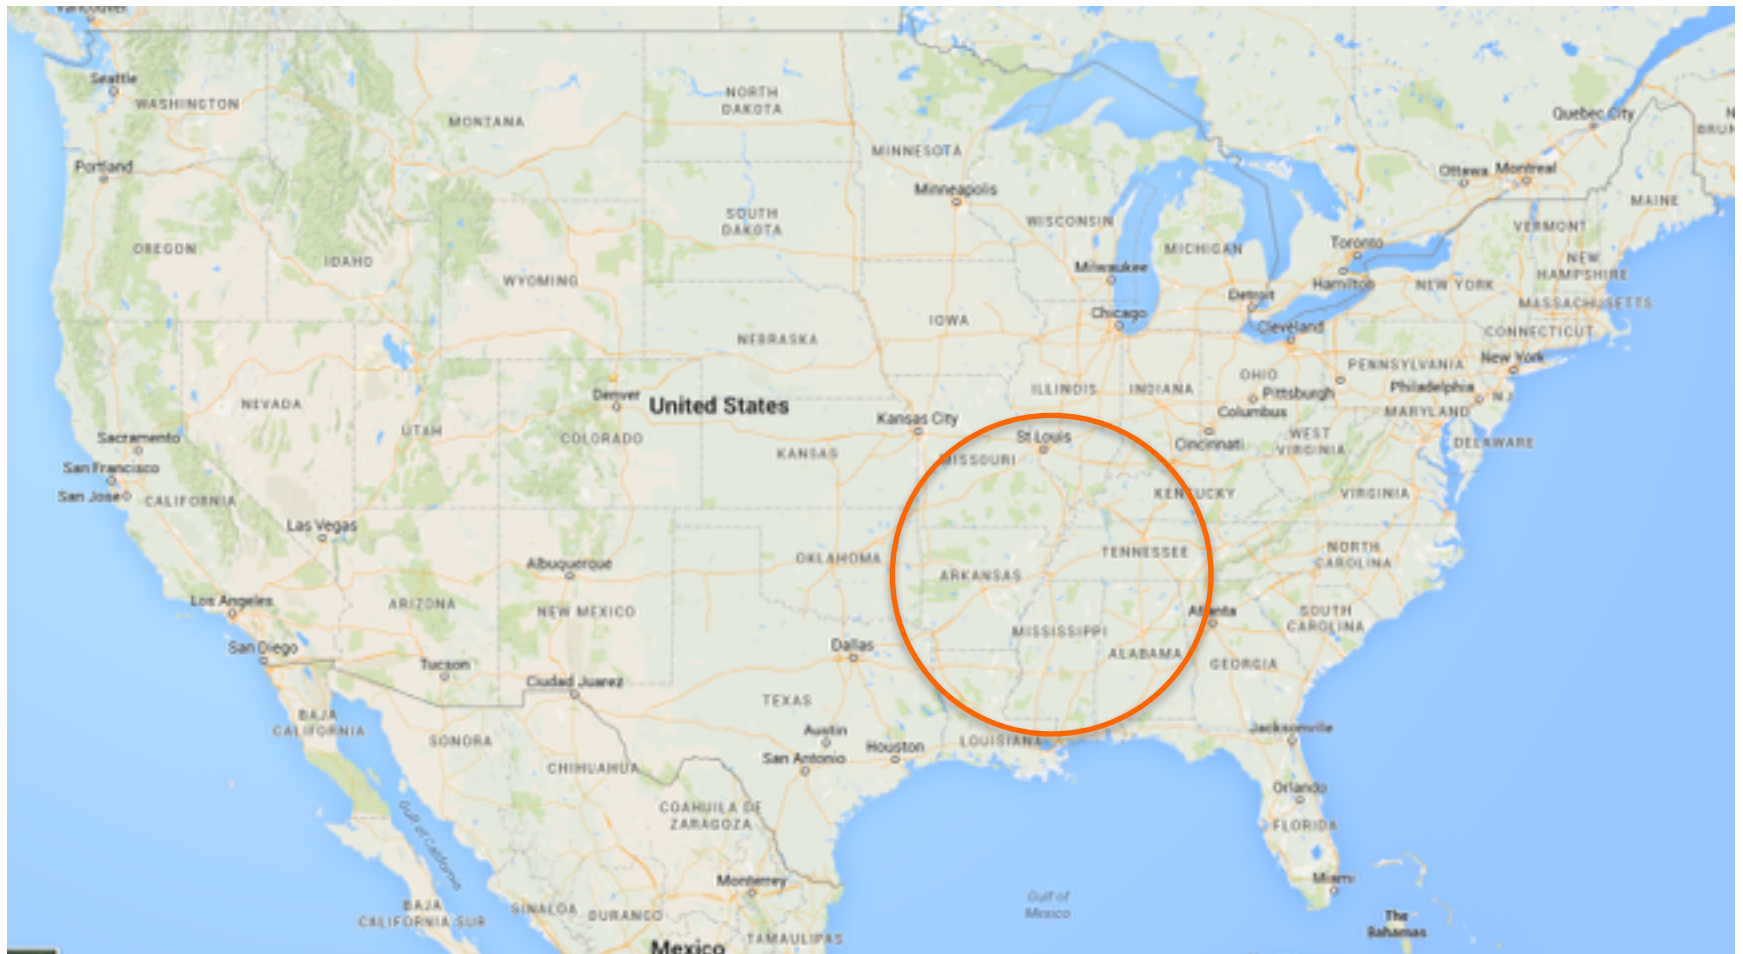
```
{ $project:
  { "_id" : 0,
    "pop1990" : 0,
    "pop2010" : 0,
    "name" : "$_id",
    "delta" :
      { "$subtract" :
        [ "$pop2010",
          "$pop1990" ] } }
}
```



```
{
  "name" : "Virginia",
  "delta" : 1813666
}
```

```
{
  "name" : "South Dakota",
  "delta" : 118176
}
```

Compare number of people living within 500KM of Memphis, TN in 1990, 2000, 2010



Compare number of people living within 500KM of Memphis, TN in 1990, 2000, 2010

```
db.cData.aggregate([
  {$geoNear : {
    "near" : {"type" : "Point", "coordinates" : [90, 35]},
    "distanceField" : "dist.calculated",
    "maxDistance" : 500000,
    "includeLocs" : "dist.location",
    "spherical" : true  }},
  {$unwind : "$data"},
  {$group : {"_id" : "$data.year",
    "totalPop" : {"$sum" : "$data.totalPop"},
    "states" : {"$addToSet" : "$name"}}},
  {$sort : {"_id" : 1}}
])
```

\$geoNear

- Order/Filter Documents by Location
 - Requires a geospatial index
 - Output includes physical distance
 - Must be first aggregation stage



\$geoNear

```
{
  "_id" : "Virginia",
  "pop1990" : 6187358,
  "pop2010" : 8001024,
  "center" :
    { "type" : "Point",
      "coordinates" :
        [78.6, 37.5]}
}
```



```
{ $geoNear : {
  "near": { "type": "Point",
            "coordinates":
              [90, 35]},
  maxDistance : 500000,
  spherical : true  }}
```



```
{
  "_id" : "Tennessee",
  "pop1990" : 4877185,
  "pop2010" : 6346105,
  "center" :
    { "type" : "Point",
      "coordinates" :
        [86.6, 37.8]}
}
```

```
{
  "_id" : "Tennessee",
  "pop1990" : 4877185,
  "pop2010" : 6346105,
  "center" :
    { "type" : "Point",
      "coordinates" :
        [86.6, 37.8]}
}
```

What if I want to save the results to a collection?

```
db.cData.aggregate([
  {$geoNear : {
    "near" : {"type" : "Point", "coordinates" : [90, 35]},
    "distanceField" : "dist.calculated",
    "maxDistance" : 500000,
    "includeLocs" : "dist.location",
    "spherical" : true  }},
  {$unwind : "$data"},
  {$group : {"_id" : "$data.year",
    "totalPop" : {"$sum" : "$data.totalPop"},
    "states" : {"$addToSet" : "$name"}}},
  {$sort : {"_id" : 1}},
  {$out : "peopleNearMemphis"}
])
```

\$out

```
db.cData.aggregate([<pipeline stages>,  
                    {"$out" : "resultsCollection"}])
```

- Save aggregation results to a new collection
- New aggregation uses:
 - Transform documents - ETL

Back To The Original Question

- Which US Division has the fastest growing population density?
 - We only want to include states with more than 1M people
 - We only want to include divisions larger than 100K square miles

Division with Fastest Growing Pop Density

```
db.cData.aggregate(  
  [{ $match : { "data.totalPop" : { "$gt" : 1000000 } } },  
    { $unwind : "$data" },  
    { $sort : { "data.year" : 1 } },  
    { $group : { "_id" : "$name",  
                  "pop1990" : { "$first" : "$data.totalPop" },  
                  "pop2010" : { "$last" : "$data.totalPop" },  
                  "areaM" : { "$first" : "$areaM" },  
                  "division" : { "$first" : "$division" } } },  
    { $group : { "_id" : "$division",  
                  "totalPop1990" : { "$sum" : "$pop1990" },  
                  "totalPop2010" : { "$sum" : "$pop2010" },  
                  "totalAreaM" : { "$sum" : "$areaM" } } },  
    { $match : { "totalAreaM" : { "$gt" : 100000 } } },  
    { $project : { "_id" : 0,  
                    "division" : "$_id",  
                    "density1990" : { "$divide" : [ "$totalPop1990", "$totalAreaM" ] },  
                    "density2010" : { "$divide" : [ "$totalPop2010", "$totalAreaM" ] },  
                    "denDelta" : { "$subtract" : [ { "$divide" : [ "$totalPop2010",  
                                                                    "$totalAreaM" ] },  
                                                       { "$divide" : [ "$totalPop1990",  
                                                                    "$totalAreaM" ] } ] } },  
                    "totalAreaM" : 1,  
                    "totalPop1990" : 1,  
                    "totalPop2010" : 1 } },  
    { $sort : { "denDelta" : -1 } } ] )
```

Aggregate Options

Aggregate options

```
db.cData.aggregate([<pipeline stages>],  
                    {'explain' : false  
                     'allowDiskUse' : true,  
                     'cursor' : {'batchSize' : 5}})
```

explain – similar to find().explain()

allowDiskUse – enable use of disk to store intermediate results

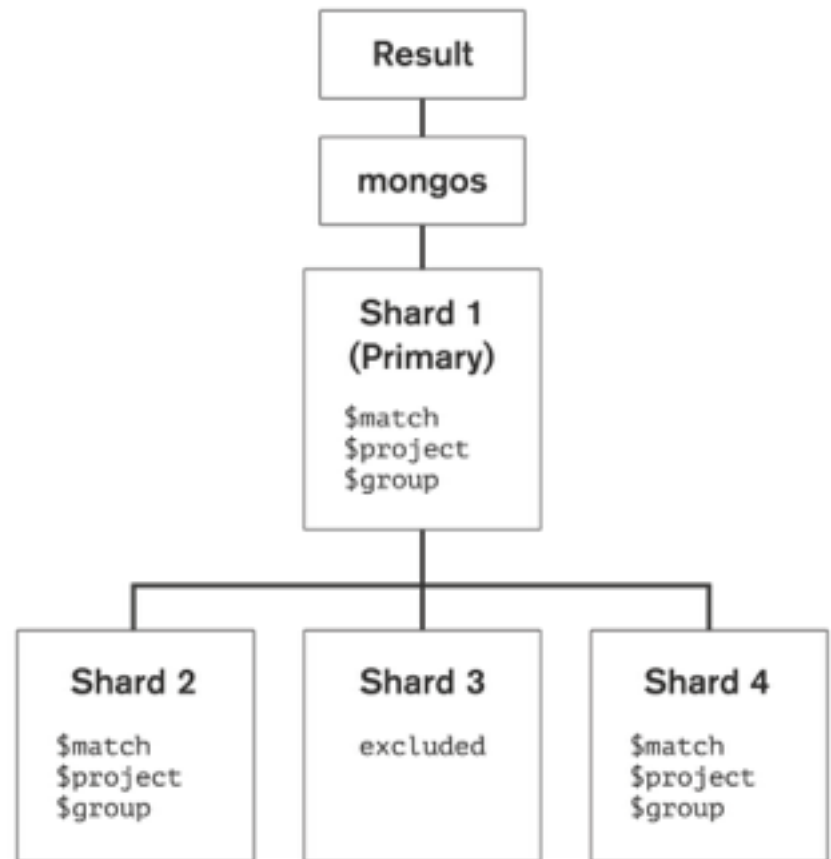
cursor – specify the size of the initial result

Aggregation and Sharding

Sharding

- Workload split between shards
 - Shards execute pipeline up to a point
 - Primary shard merges cursors and continues processing*
 - Use explain to analyze pipeline split
 - Early `$match` may excuse shards
 - Potential CPU and memory implications for primary shard host

*Prior to v2.6 second stage pipeline processing was done by mongos



Summary

Analytics in MongoDB?

Create
Read
Update
Delete



YES!

?



Analytics



Group
Count
Derive Values
Filter
Average
Sort

Framework Use Cases

- Basic aggregation queries
- Ad-hoc reporting
- Real-time analytics
- Visualizing and reshaping data

Questions?

shannon@mongodb.com